

# Introduction to



By

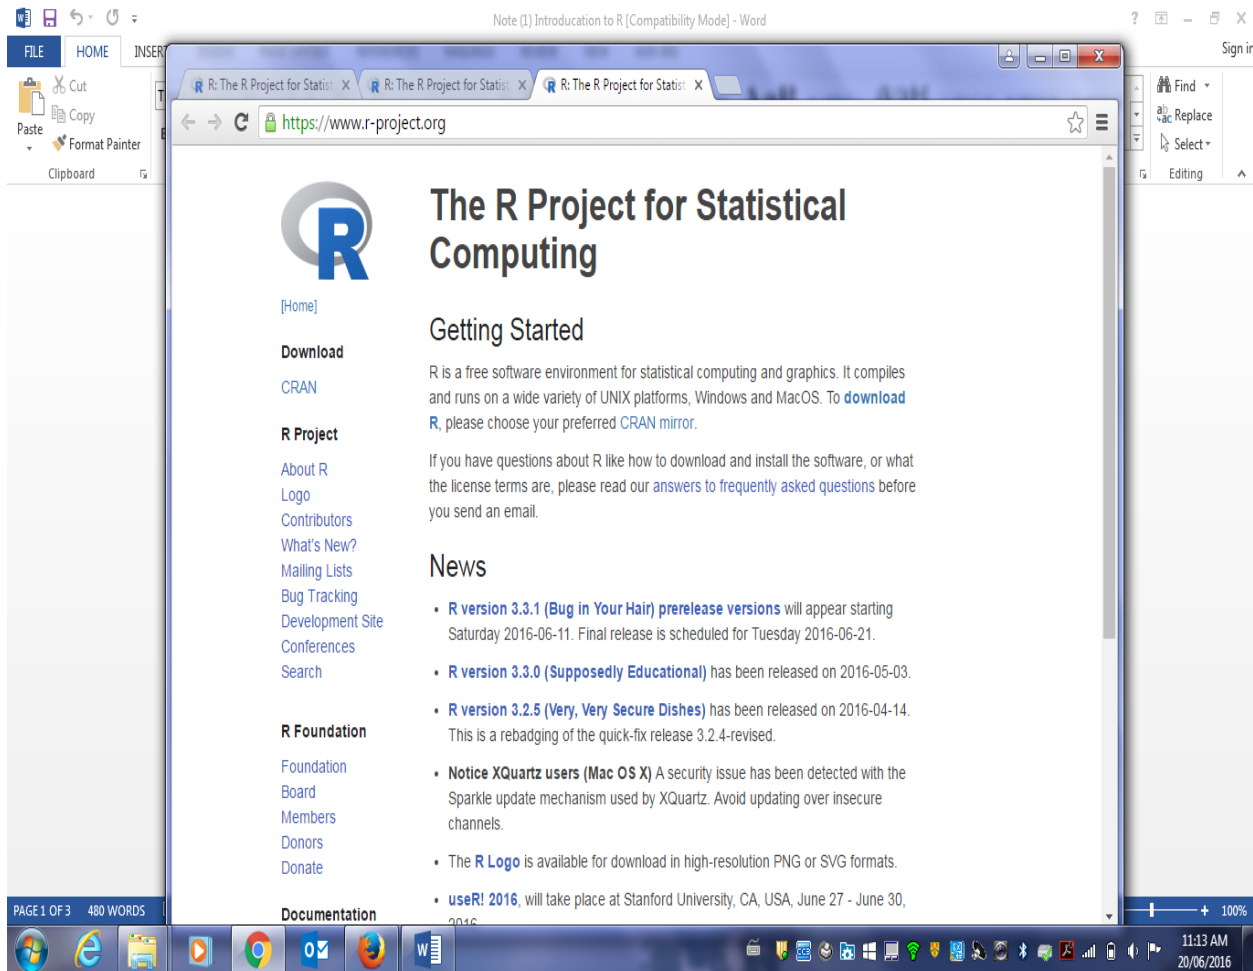
Dina Dawoud

and

Nagham Mohammad

# What is R?


According to their site <http://www.r-project.org/> :



"R is a language and environment for statistical computing and graphics."

"R provides a wide variety of statistical (linear and nonlinear modelling, classical statistical tests, time-series analysis, classification, clustering ...) and graphical techniques, and is highly extensible."

R was created by [Ross Ihaka](#) and [Robert Gentleman](#) at the [University of Auckland](#), New Zealand, and is currently developed by the *R Development Core Team*.



Robert Gentleman



Ross Ihaka



Photo: <http://www.nytimes.com/2009/01/07/technology/business-computing/07program.html>

## Obtaining a copy of the R applications

R is free software - see the R site above for the terms of use. It runs on a wide variety of platforms including UNIX, Windows and MacOS.

Download a copy of this application from their site: <http://www.r-project.org/>

The screenshot shows a web browser window displaying the homepage of The R Project for Statistical Computing. The browser's address bar shows the URL <https://www.r-project.org>. The page features the R logo, a navigation menu on the left, and a main content area with sections for 'Getting Started', 'News', and 'Documentation'. The 'Getting Started' section includes a link to download R and choose a CRAN mirror. The 'News' section lists recent updates, including R version 3.3.1, 3.3.0, and 3.2.5, as well as a security issue for XQuartz users on Mac OS X. The 'Documentation' section is partially visible at the bottom. The browser window is titled 'R: The R Project for Statistical Computing'. The background shows a Windows desktop with a Word document open, titled 'Note (1) Introduction to R [Compatibility Mode] - Word'. The taskbar at the bottom displays various application icons, including the Start button, Internet Explorer, File Explorer, and several instances of the R application. The system clock in the bottom right corner shows the time as 11:13 AM on 20/06/2016.

Page 1 of 3 480 WORDS

## **R**

Organization Web Site: <http://www.r-project.org/>

Any Special Technical Requirements: This is an open source product available for a variety of computing platforms.

[R for Mac OS X FAQ](#)

[R for Windows FAQ](#)

Available online:

## **Helpful Resources**

The following references would be a good step towards building a solid foundation in using R.

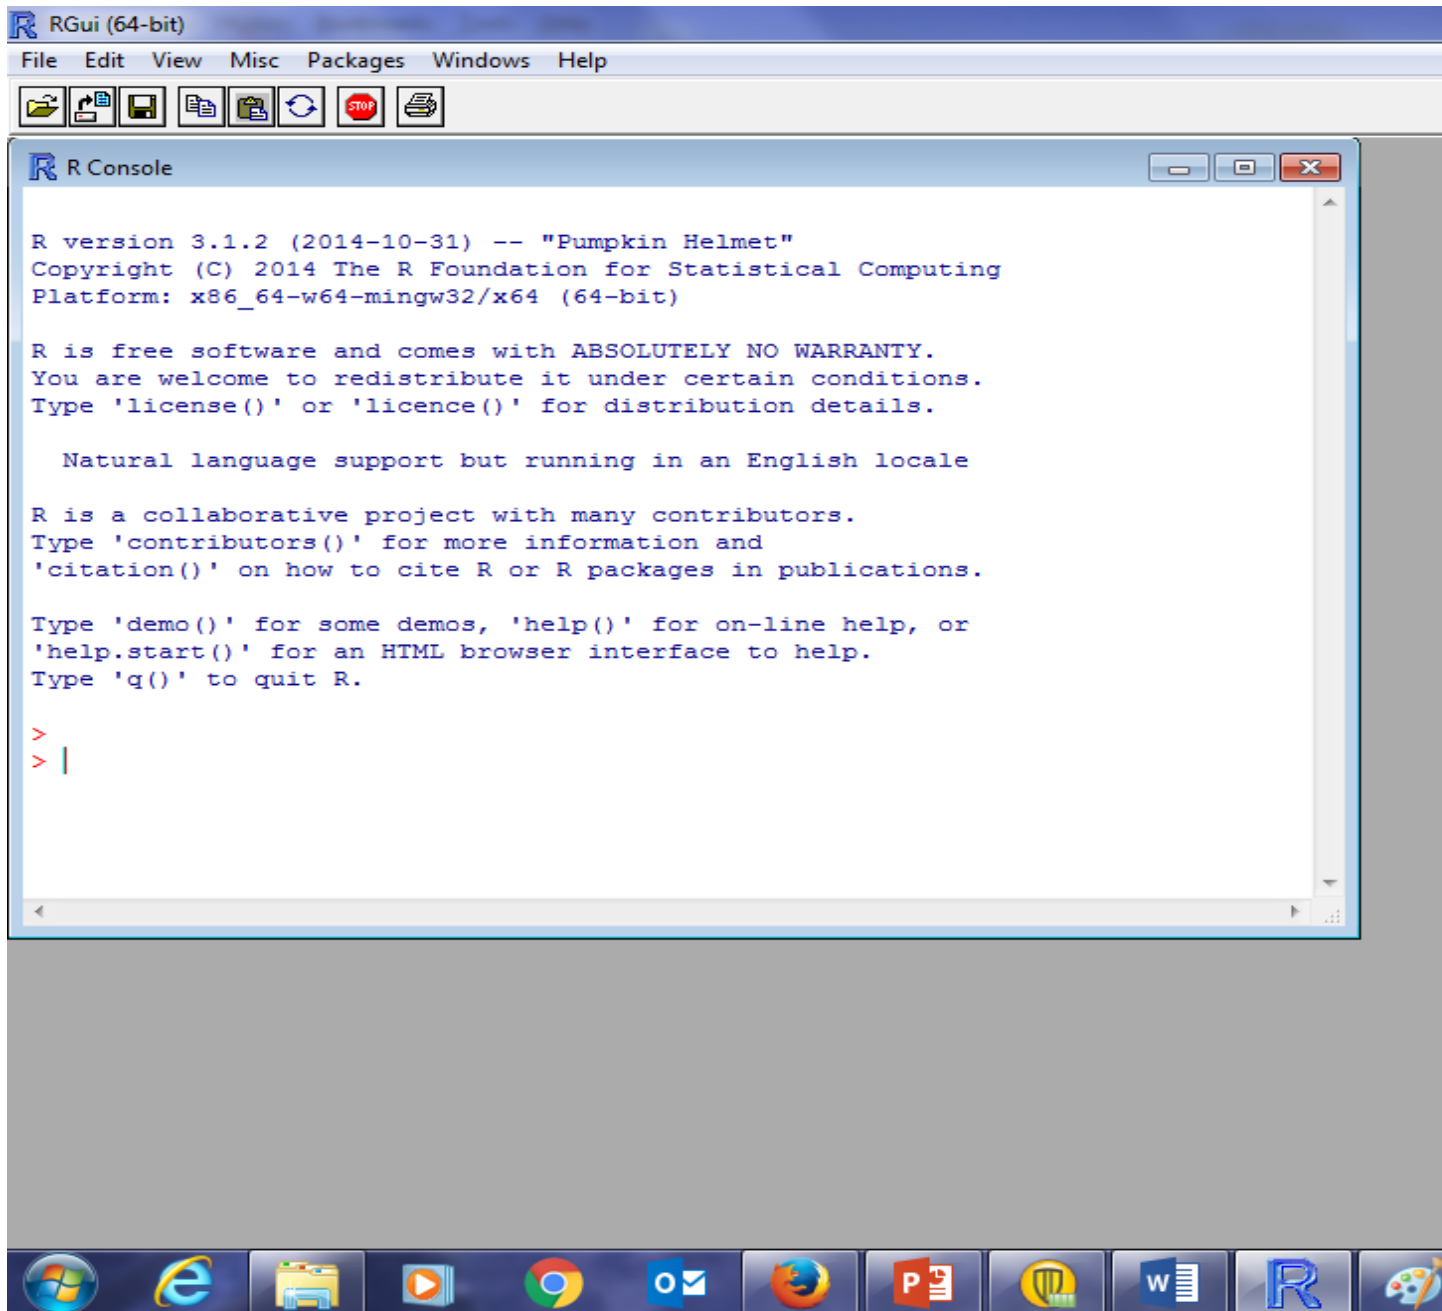
- The [R Project Homepage](#)
- [An Introduction to R](#) - written by W. N. Venables, D. M. Smith and the R Development Core Team
- [Short Reference Card](#) - written by Tom Short, EPRI PEAC
- [R Seek](#) helps you find the R function you require

## **Book**

Modern Applied Statistics with S. Fourth Edition by W. N. Venables and B. D. Ripley

## Starting and Quitting R.

- Double clicking on the R icon starts the program.
- The first thing that will happen is that R will open the *console*, into which the user can type commands.



- The greater-than sign (>) is the prompt symbol. When this appears, you can begin typing commands.
- R can be used as a calculator.
- We can type simple arithmetical expressions at the > prompt:

15 + 10

- Upon pressing the Enter key, the result 25 appears, prefixed by the number 1 in square brackets:

[1] 25

- The [1] indicates that this is the first (and in this case only) result from the command.

## Notes

### 1) Help

If you need help understanding a command or its syntax type either **(?) command**, or **(help) command** and R will display the help available on this topic.

Examples

?mean

help(var)

2) Certain functions in R may NOT run on all platforms (e.g., Windows, MAC, Linux, etc...) the same way.

# Comments

Words after a pound sign (#) are considered comments and will not be executed by R.

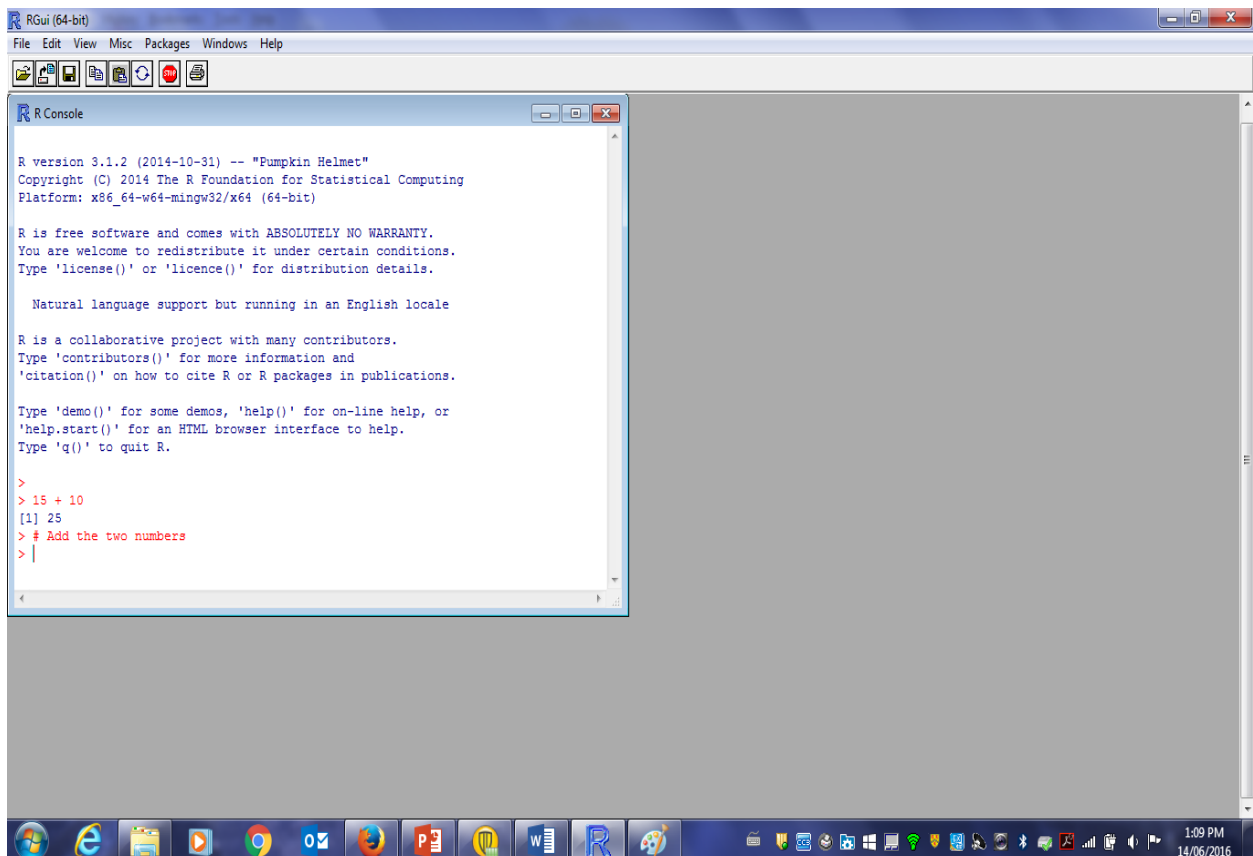
## Examples

```
2 * 5 # "*" is the symbol for multiplication.  
[1] 10
```

```
3 - 10  
[1] -7
```

```
12 / 4  
[1] 3
```

```
3^4  
[1] 81
```



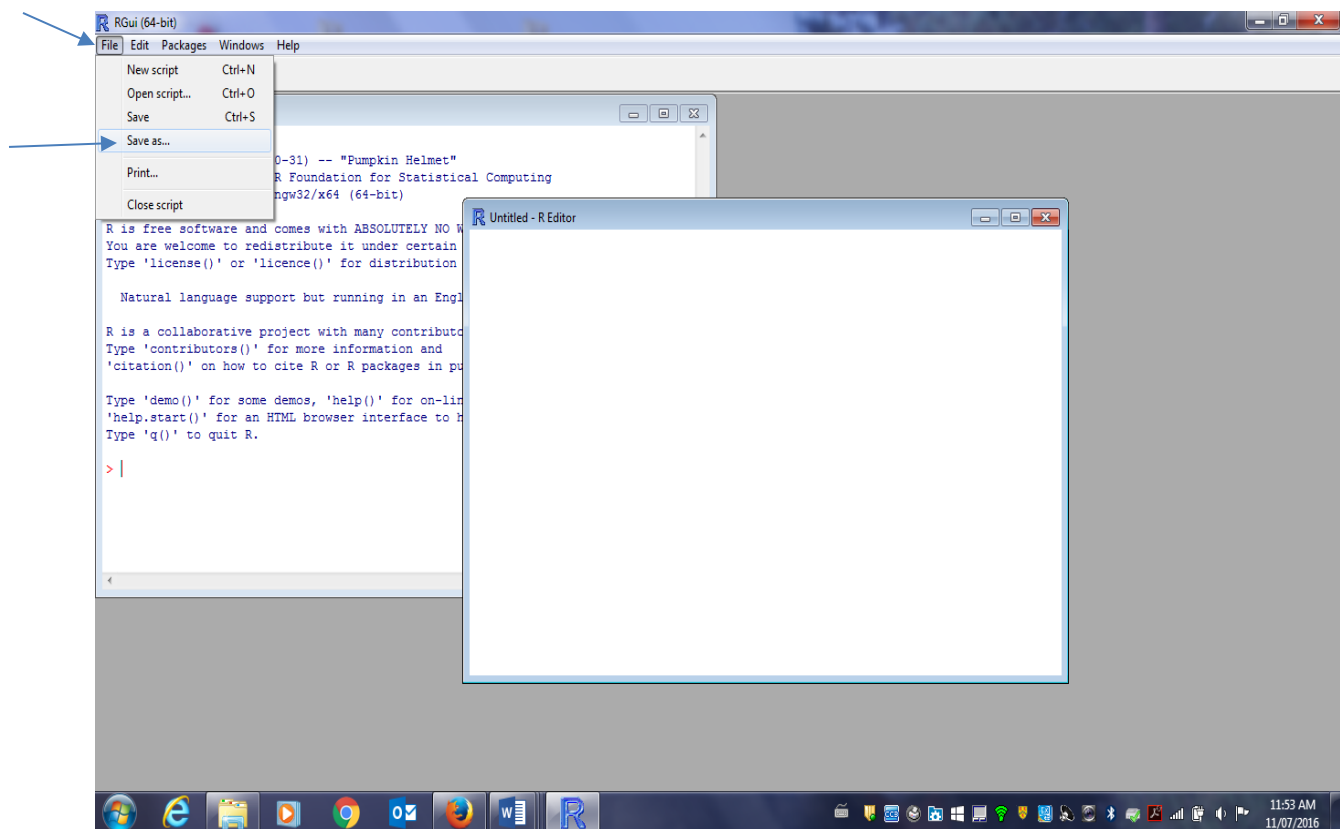


- To quit your R session, type  
`q()`

If you then hit the Enter key, you will be asked whether to save an image of the current workspace, or not, or to cancel. The workspace image contains a record of the computations you've done, and may contain some saved results. Hitting the Cancel option allows you to continue your current R session. We rarely save the current workspace image, but occasionally find it convenient to do so.

## Recording your work

Rather than saving the workspace, we prefer to save the commands we entered. In Windows, the easiest way to do this is from the File menu. Open new script. At the end of a session, save the final script for a permanent record of your work.



# Named storage

R has a workspace known as the *global environment* that can be used to store the results of calculations, and many other types of objects.

## Example

- Suppose we would like to store the result of the calculation  $5^5$  for future use.
- We will assign this value to an object called X. To this, we type

```
X <- 5^5
```

- We tell R to make the assignment using an arrow that points to the left, created with the less-than sign (<) and the hyphen (-).
- R also supports using the equals sign (=) in place of the arrow in most circumstances, but we recommend using the arrow, as it makes clear that we are **requesting an action** (i.e. an assignment), rather than **stating a relation** (i.e. that X is equal to  $5^5$ ).
- Note that when we hit Enter, nothing appears on the screen except a new prompt: R has done what we asked, and is waiting for us to ask for something else.
- We can see the results of this assignment by typing the name of our new object at the prompt:

```
X  
[1] 3125
```

## Installing R packages

Part of the reason R has become so popular is the vast array of packages available at the [cran](#) and [bioconductor](#) repositories.

Go to the packages menu item ...Click ... go to install package(s)...You will see a list of CRAN mirror where you can download ... Click on the mirror close to you (Canada ON)... R will take a couple of moments and then will provide you with all the available packages...Click on the one you need or a list of them by shift and click on others items.

If you want to choose more than one and they are not on order

Click on Control and hold it until you choose all of the packages you need.

Then click on ok

If you need more information about the packages in R go to the CRAN home page

<http://cran.r-project.org/>

To the left hand side go down to the link of packages click on that and it will take you to the contributed packages in R

## Available Packages

Currently, the CRAN package repository features 7415 available packages.

[Table of available packages, sorted by date of publication](#)

[Table of available packages, sorted by name](#)

You will see all the available packages in R the same ones at the packages menu box. Or, you can install packages using the command

```
install.packages("ggplot2")
```

```
library(ggplot2)
```

# Functions

Most of the work in R is done through *functions*.

Example (1)

Consider the standard formula for the sample mean of a sample  $x_1, \dots, x_n$  is

$$\bar{x} = \frac{\sum_{i=1}^n x_i}{n}$$

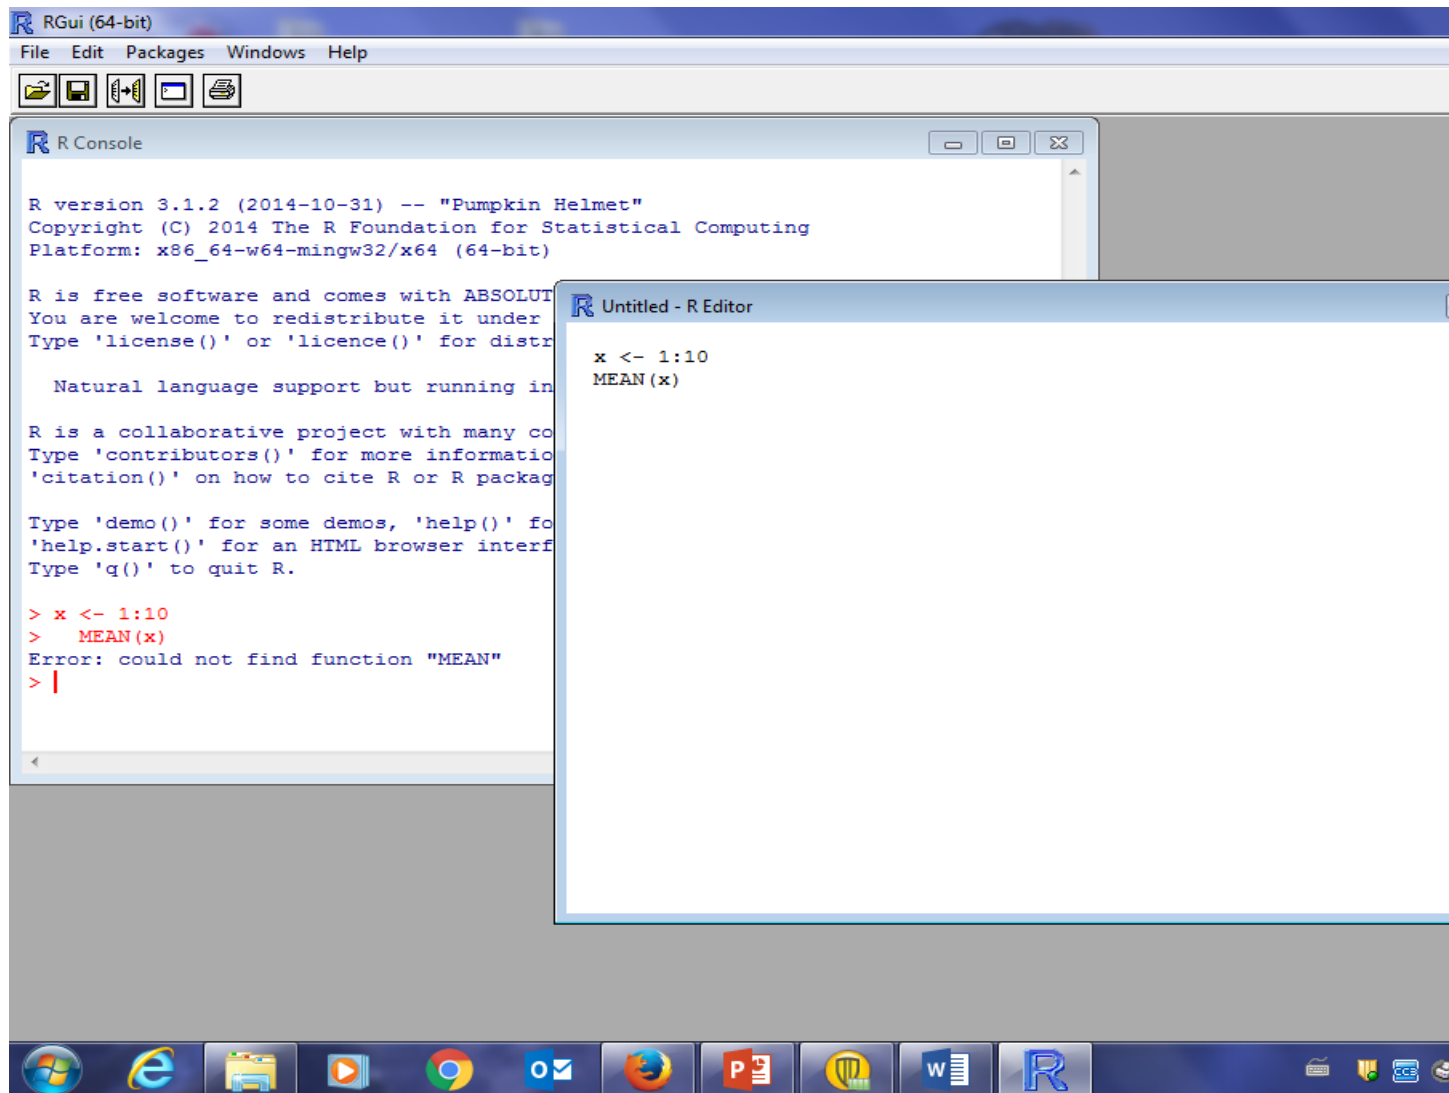
In R, mean (). For example:

```
x <- 1:11  
  
mean(x)  
[1] 6
```

## R is case-sensitive

See what happens if you type

```
x <- 1:10  
MEAN(x)  
Error: could not find function "MEAN"
```



Or

```
mean(x)
[1] 5.5
```

Now try

```
MEAN <- mean
MEAN(x)
[1] 5.5
```

The function `mean()` is built in to R. R considers `MEAN` to be a different function, because it is case-sensitive: `m` is a different letter than `M`.

## Vectors

- A numeric vector is a list of numbers. The `c( )` function is used to collect things together into a vector. We can type

```
c(0, 7, 8)
[1] 0 7 8
```

- We can assign this to a named object:

```
x <- c(0, 7, 8) # now x is a 3-element vector
```

To see the contents of `x`, simply type

```
x
[1] 0 7 8
```

- `Symbol` can be used to create sequences of increasing (or decreasing) values.

Example

```
numbers2to30 <- 2:30
numbers2to30
```

```
numbers2to30
[1] 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25
[25] 26 27 28 29 30
```

```
x <- 1:11
x
```

```
[1] 1 2 3 4 5 6 7 8 9 10 11
```

- Vectors can be joined together (i.e. concatenated) with the `c` function.

Example

```
newvector <- c(numbers2to30, x)
newvector
```

```
[1] 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25
[25] 26 27 28 29 30 1 2 3 4 5 6 7 8 9 10 11
```

## Extracting elements from vectors

- A nicer way to display the 22nd element of `newvector` is to use square brackets to extract just that element:

```
newvector[22]
[1] 23
```

- To get the fourth through eighth elements

```
newvector[4:8]
[1] 5 6 7 8 9
```

- Negative indices can be used to avoid certain elements. For example, we can select all but the second element of `newvector` as follows:

```
newvector [-2]

[1] 2 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26
[25] 27 28 29 30 1 2 3 4 5 6 7 8 9 10 11
```

- The third through eleventh elements of `newvector` can be avoided as follows:

```
newvector [-(3:11)]

[1] 2 3 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 1 2 3 4
[25] 5 6 7 8 9 10 11
```

- Do not mix positive and negative indices. To see what happens, consider

```
newvector [c(-2, 3)]
```

Error in newvector[c(-2, 3)] :

The problem is that it is not clear what is to be extracted: do we want the third element of **newvector** before or after removing the second one?

## Vector Arithmetic

- Arithmetic can be done on R vectors. For example, we can multiply all elements of x by 3:

```
x * 3
[1] 3 6 9 12 15 18 21 24 27 30 33
```

Note that the computation is performed element wise.

- Addition, subtraction and division by a constant have the same kind of effect. For example,

```
y <- x - 5
y
[1] -4 -3 -2 -1 0 1 2 3 4 5 6
```



## Simple Patterned Vectors

- Patterned vectors can be produced using the `seq()` function as well as the `rep()` function.

Example, the sequence of odd numbers less than or equal to 17 can be obtained using

```
seq(1, 17, by=2)
[1] 1 3 5 7 9 11 13 15 17
```

Notice the use of `by=2` here. The `seq()` function has several *optional parameters*, including one named `by`. If `by` is not specified, the default value of 1 will be used.

- Repeated patterns are obtained using `rep()`. Consider the following examples:

```
rep(9, 12) # repeat the value 9, 12 times
[1] 9 9 9 9 9 9 9 9 9 9 9 9
```

```
rep(seq(2, 20, by=2), 2) # repeat the pattern 2 4 ... twice
```

```
[1] 2 4 6 8 10 12 14 16 18 20 2 4 6 8 10 12 14 16 18 20
```

## Data frames (data.frame)

Data sets frequently consist of more than one column of data, where each column represents measurements of a single variable. Each row usually represents a single observation. This format is referred to as *case-by-variable* format.

### Example

The following data set consists of four observations on the three variables x, y, z

x	y	z
61	13	4
175	21	18
111	24	14
124	23	18

```
x <- c( 61,175,111,124)
y <- c( 13,21,24,23)
z <- c( 4,18,14,18)
Mydata <- data.frame( x,y,z)
Mydata
```

You can see the names of the objects in your data.frame using the names()function,

```
names( Mydata)
```

You can extract parts of your data.frame

```
Mydata$x
Mydata$y
Mydata$z
```

## Example

```
colors <- c("red", "yellow", "blue")
numbers <- c(1, 2, 3)
Data1 <- data.frame(colors, numbers)
Data1
```

```
  colors numbers
1   red      1
2 yellow     2
3  blue      3
```

## Practice Question

Using `rep()` and `seq()` as needed, create the vectors

(a) 0 0 0 0 0 1 1 1 1 1 2 2 2 2 2 3 3 3 3 3 4 4 4 4 4

and

(b) 1 2 3 4 5 1 2 3 4 5 1 2 3 4 5 1 2 3 4 5

## Solution

```
(a) X <- seq(0, 4, by=1)    ## or you can X<- c( 0, seq(1:4))
rep( X, each=5)
```

```
[1] 0 0 0 0 0 1 1 1 1 1 2 2 2 2 2 3 3 3 3 3 4 4 4 4 4
```

```
(b) rep( seq(1:5), 4)
```

```
[1] 1 2 3 4 5 1 2 3 4 5 1 2 3 4 5 1 2 3 4 5
```

## Practice Problem

Consider the pressure data frame in R. There are two columns: temperature and pressure.

### **To work with the data set `pressure` in R**

Step (1)

Type `pressure` in R, you suppose to get the data set, if not probably you need to install some packages and R will instruct you to their names.

packages ----- Install Package(s)--- Canada(ON)----- the name of the package

Step(2)

library(the name of the package you install)

Step(3)

Now type `pressure` again you should get the data

Step(4)

To obtain a specific column in pressure, use:

```
x1<- pressure$temperature
```

```
x2<-pressure$pressure
```

**Note : Before working with the data do the following**

Add two numbers to each observation of the variable that you will work with based on your student number; so if my student number is (117893) I will add the number 93.

### Example

Suppose the first observation (data point) in your variable is 7 then it will be  $(7+93=100)$  and the same for all the other  $n$  remaining observations.

Note: Do not do that by your hand let R do that for you, use

```
x1<- x1+93  
x2<-x2+93.
```

- (a) Find the mean, median, variance, and standard deviation of the temperature variable.
- (b) Find the mean, median, variance, and standard deviation of the pressure variable.
- (c) Plot the probability histogram for the temperature and the pressure using the par function to display them in a  $2 \times 1$  layout on the graphics page. Repeat once again using a  $1 \times 2$  layout with different colors, use

```
par(mfrow=c(2,1))  
par(mfrow=c(1,2))
```

Add a suitable title and labels to the graphs.

## Solution

pressure

```
x1 <- pressure$temperature
```

```
x2 <- pressure$pressure
```

```
x1 <- x1+93 #temperature
```

```
x2 <- x2+93 #pressure
```

```
mean(x1) # [1] 273
```

```
median(x1) #273
```

```
sd(x1) #[1] 112.5463
```

```
par(mfrow=c(2,1))
```

```
hist(x1,col="blue",probability = TRUE,ylab="probability",  
xlab="temperature",main="Histogram of the temperature variable ")
```

```
hist(x2,col="red",probability = TRUE,ylab="probability",  
xlab="pressure",main="Histogram of the pressure variable")
```

```
par(mfrow=c(1,2))
```

```
hist(x1,col="blue",probability = TRUE,ylab="probability",  
xlab="temperature",main="Temperature variable ")
```

```
hist(x2,col="red",probability = TRUE,ylab="probability",  
xlab="pressure",main="Pressure variable")
```

